


Cryptocurrency Market Sentiment Analysis: A Transformers, Lexicon and classical Models based approach

Teddy Thomas



Motivation

Crypto market is highly sentiment-driven

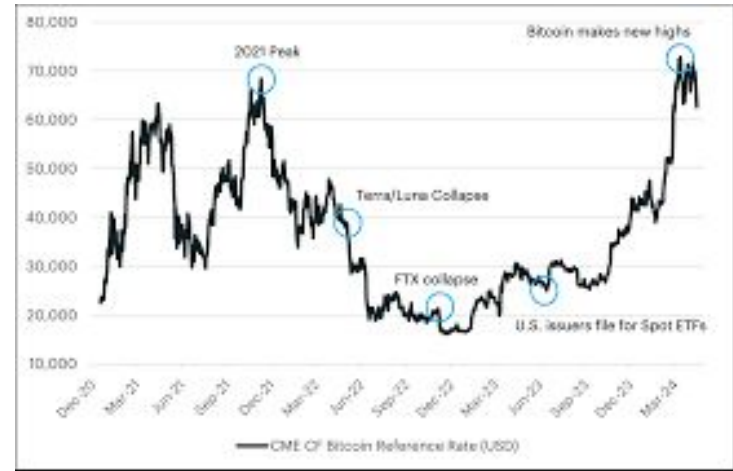
Retail investors rely on platforms like **StockTwits**

Predicting sentiment helps with:

- Automated trading systems

- Market trend analysis

- Retail investor sentiment tracking



Problem Statement

?

How can we accurately classify the sentiment
(**Bullish, Bearish, or Neutral**) of user-generated posts
on **StockTwits** related to **cryptocurrencies** such as
Bitcoin, Ethereum, and Shiba Inu using a
combination of **traditional machine learning**,
lexicon-based, and **transformer-based NLP models**?

\$BTC to the moon 🚀!

Objective

Classify crypto-related StockTwits posts into:

Bullish (2)

Bearish (1)

Neutral (0)

Compare performance of:

Lexicon-based models

Traditional ML classifiers

Word embeddings

Transformer-based models

Dataset Description: ElKulako/stocktwits-crypto

Platform: [StockTwits](#)

Collection Period:

Training: 1 Nov 2021 – 15 Jun 2022

Testing: 16 Jun 2022 – 30 Jun 2022

Volume:

Total StockTwits Posts: 1.875 million

Training Set: 1.332 million posts

Test Set: 83,257 posts

Sentiment Labels:

Bullish / Bearish (self-labeled by users)

Neutral inferred when no label is given

Cryptocurrencies Covered:

Bitcoin (BTC.X)

Ethereum (ETH.X)

Shiba Inu (SHIB.X)

(Only these are used for supervised training)

Preprocessing Steps:

Removed: URLs, usernames (@), cashtags (\$), hashtags (#), wallet addresses, emojis (except for LUKE use), non-English scripts

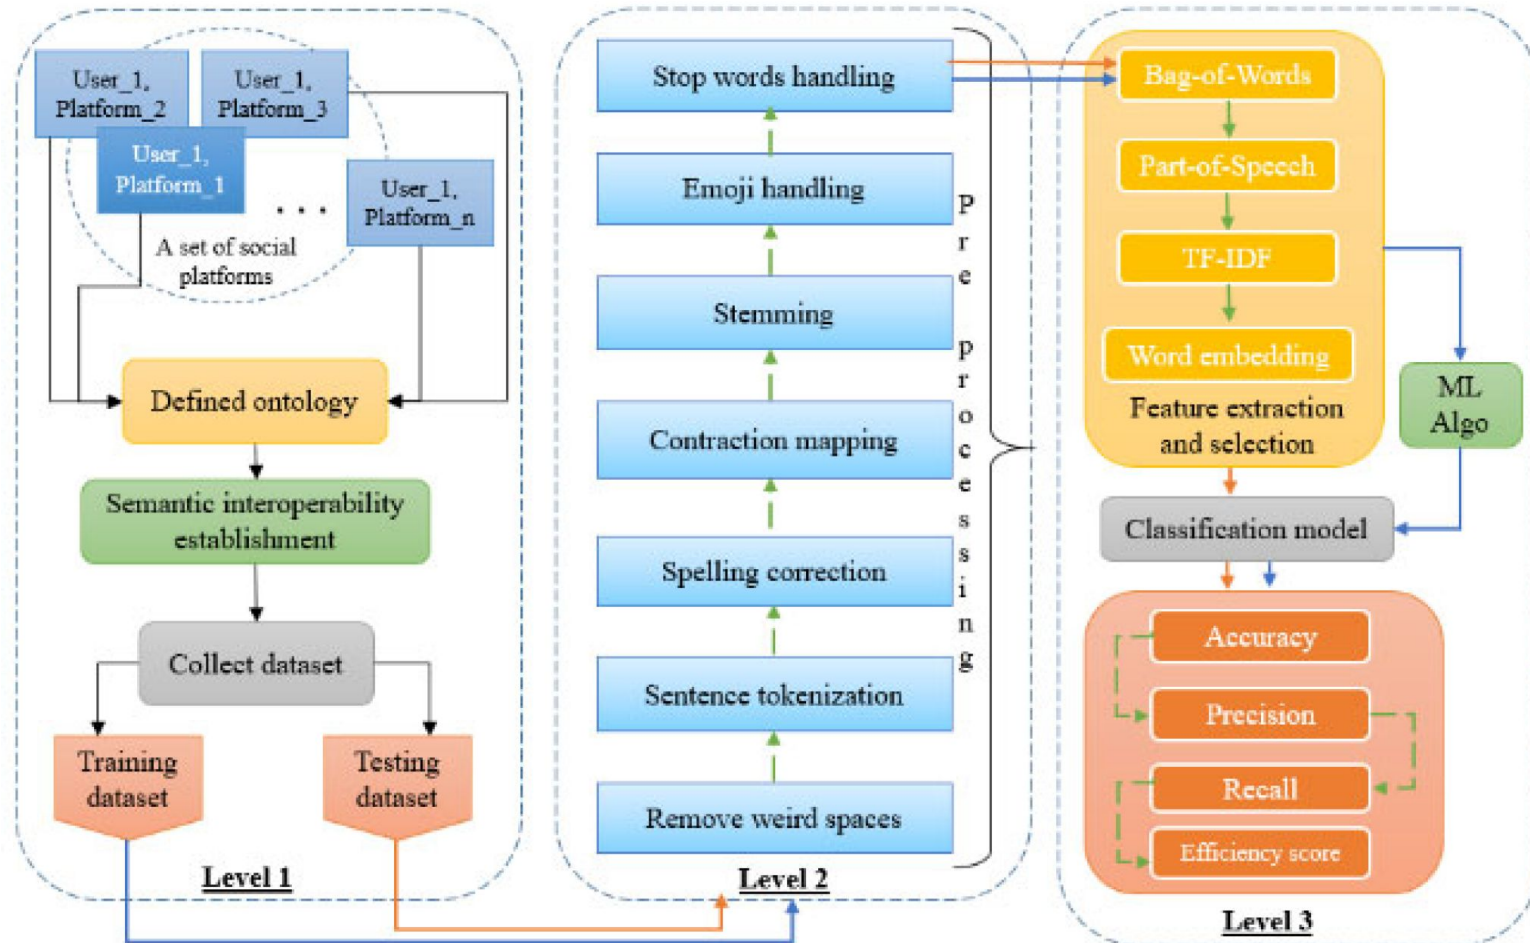
Fixed: encoding issues, multiple dots/spaces

Lowercased text, dropped duplicates and short posts (<4 words)

@elonmusk \$BTC 🤖🤖💥💥

->just BTC

Sentiment Analysis: Methodology



Methodology Overview

Lexicon-Based: VADER SentimentIntensityAnalyzer

Classical ML:

CountVectorizer + Naive Bayes

CountVectorizer + Logistic Regression




TF-IDF + Linear SVM

Embeddings:

Word2Vec + Logistic Regression

Transformer-Based:


BERT Sentence Embeddings + ClassifierDL (Spark NLP)





Choosing What to Use

Hypothesis



Goal

Best Choice

Simple model + small data

CountVectorizer or TF-IDF

Classical ML + interpretability

TF-IDF

Deep learning

Word2Vec / GloVe / FastText

Best accuracy

BERT embeddings / fine-tuning


Sentence-level meaning

Sentence Transformers / USE

Real-time / speed-sensitive

Hashing Vectorizer or distilled models

Exploratory Data Analysis



Label distribution:

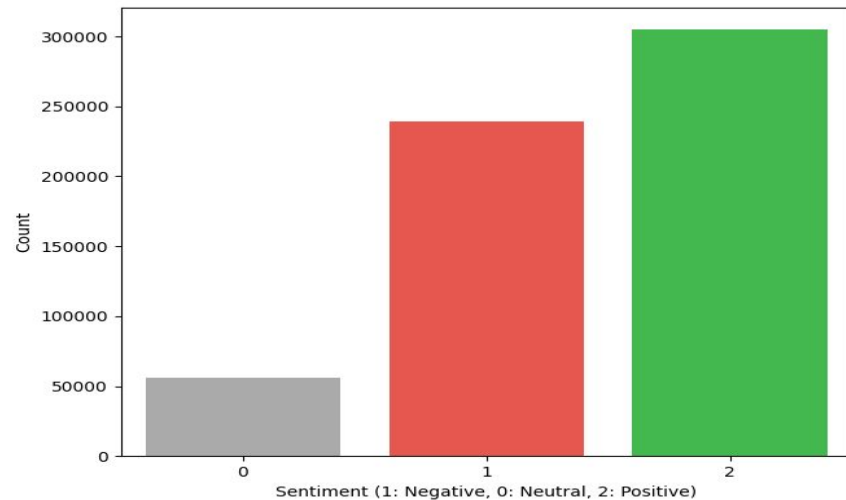
sentiment

Bullish 305092


Neutral 239054

Bearish 55854


Sentiment Distribution




Text Length Distribution by Sentiment




Word Count Distribution by Sentiment



Average Word Length by Sentiment



Correlation Between Text Features and Sentiment



Big Data & Scalability


Apache Spark on Databricks: Used for distributed processing of 1.8M+ posts using Spark DataFrames, MLlib pipelines, and UDFs.

Parallelism: Tokenization, vectorization (TF-IDF, Word2Vec), and model training were executed in parallel across Spark workers for scalability.

Batch Processing: Data was processed in partitions to handle memory efficiently, especially during transformation and model inference stages.

CBCB HPC Cluster: Employed for heavy transformer-based models like BERT; used SLURM for job scheduling and high-performance GPU computation.

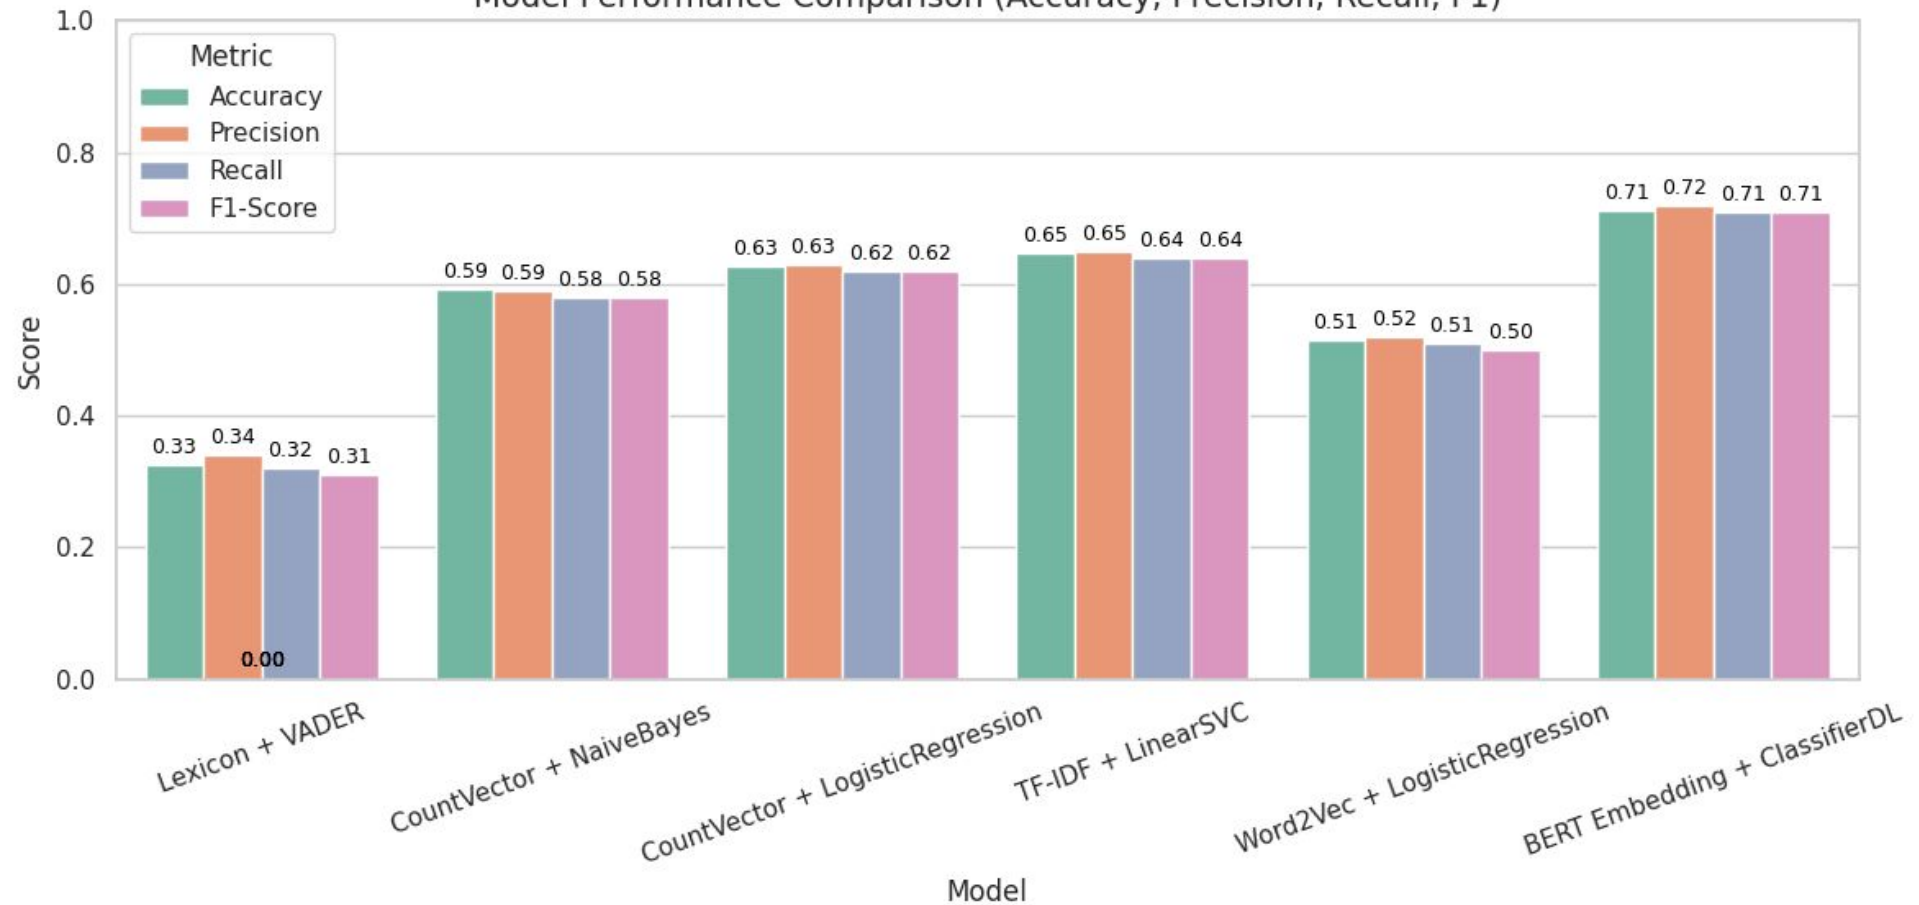
Hybrid Architecture: Combined cloud (Databricks) and HPC (CBCB) for optimal resource usage and faster execution of large-scale NLP tasks.



Model Evaluation

Model	Accuracy	Precision	Recall	F1-Score
VADER	0.3253	0.34	0.32	0.31
Naive Bayes	0.5911	0.59	0.58	0.58
Logistic Regression	0.6267	0.63	0.62	0.62
Linear SVC	0.6472	0.65	0.64	0.64
Word2Vec + LR	0.5145	0.52	0.51	0.50
BERT (Spark NLP)	0.7118	0.72	0.71	0.71

Model Performance Comparison (Accuracy, Precision, Recall, F1)



Analysis and Insights

Lexicon-based approach underperforms on short, informal text

Classical models improve with better vectorization

Transformer-based model performed best

Word2Vec was sensitive to training data quality

Trade-off: Simpler models are faster, but less accurate



Challenges

Handling imbalanced and noisy sentiment labels

Short length of social media posts

Sparse text with emojis, tickers, slang

Limited labeled neutral examples (inferred)

References & Citations

- [1] Dataset: M. Kulakowski and F. Frasinicar, "Sentiment Classification of Cryptocurrency-Related Social Media Posts," in IEEE Intelligent Systems, vol. 38, no. 4, pp. 5-9, July-Aug. 2023, doi: 10.1109/MIS.2023.3283170.
- [2] J. Devlin et al., "BERT: Pre-training of deep bidirectional transformers for language understanding," in Proc. Conf. North Amer. Chapter Assoc. Comput. Linguistics, Hum. Lang. Technol. (NAACL-HLT), 2019, pp. 4171–4186.
- [3] © Databricks 2025. All rights reserved. Apache, Apache Spark, Spark and the Spark logo are trademarks of the Apache Software Foundation.

Thank YOU!

What is the Ground Truth Here?

Your dataset provides sentiment labels based on **user annotations**:

- **Bullish** → Label: 2(Positive sentiment)
- **Bearish** → Label: 1X(Negative sentiment)
- **Neutral** → Label: 0 (Inferred when no explicit sentiment label is provided)

These sentiment labels are either:

- **Explicitly provided by users** (Bullish/Bearish tags on StockTwits), or
- **Inferred** (Neutral, when no tag exists).